

▼ **Black Swan Anomalies Stock Market Analysis**

**Authors**

Oluwatamilore (Tami) Ajibade, Melina Nguyen, Oluwatobi (Tobi) Oyinloye

▼ **Introduction**

```
from google.colab import files

uploaded = files.upload()
```

Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving weekly TRM.csv to weekly TRM (?).csv

```
import pandas as pd

stocks_df = pd.read_csv("weekly_IBM.csv")
stocks_df.head()
```

	timestamp	open	high	low	close	volume
0	2023-12-01	154.99	160.590	154.75	160.55	21900644
1	2023-11-24	152.51	155.705	152.35	155.18	11362696
2	2023-11-17	148.46	153.500	147.35	152.89	19547595
3	2023-11-10	147.89	149.680	145.28	149.02	18357944
4	2023-11-03	143.19	148.445	142.58	147.90	22959464

We also will define our start dates and end dates for our three anomalies.

```

## anomaly 1 - 9/11
# 4 months of 9/11 financial crisis
anom_1a_start = "2001-09-11"
anom_1a_end = "2002-01-11"

# 4 months BEFORE 9/11 financial crisis
anom_1b_start = "2001-05-10"
anom_1b_end = "2001-09-10"

# 4 months AFTER 9/11 financial crisis
anom_1c_start = "2002-01-11"
anom_1c_end = "2002-05-01"

## anomaly 2 - 2008 Housing Market Crash
# 1 year of 2008 Housing financial crisis
anom_2a_start = "2008-01-01"
anom_2a_end = "2008-12-31"

# 1 year BEFORE 2008 Housing financial crisis
anom_2b_start = "2007-01-01"
anom_2b_end = "2007-12-31"

# 1 year AFTER 2008 Housing financial crisis
anom_2c_start = "2009-01-01"
anom_2c_end = "2009-12-31"

## anomaly 3 - COVID-19 pandemic
# 2 months of COVID-19 financial crisis
anom_2a_start = "2020-03-01"
anom_2a_end = "2020-04-31"

# 2 months BEFORE COVID-19 financial crisis
anom_2b_start = "2020-01-01"
anom_2b_end = "2020-02-29"

# 2 months AFTER COVID-19 financial crisis
anom_2c_start = "2020-05-01"
anom_2c_end = "2009-07-01"

```

## Methodology and Results

### Data Preparation

We apply a log transformation to our data. Stock prices are typically positively skewed and exhibit heteroscedasticity (the spread of data points changes as the values of the independent variable change). Applying a log transformation to the stock prices can help stabilize variances and make the data more suitable for time series analysis or modeling.

```

import numpy as np
def logs_to_stocks(x):
    if pd.api.types.is_numeric_dtype(x):
        return np.log(x)
    else:
        return x

logs_df = stocks_df.apply(logs_to_stocks)
logs_df.head()

```

	timestamp	open	high	low	close	volume
0	2023-12-01	5.043361	5.078855	5.041811	5.078605	16.902027
1	2023-11-24	5.027230	5.047963	5.026181	5.044586	16.245846
2	2023-11-17	5.000316	5.033701	4.992811	5.029719	16.788363
3	2023-11-10	4.996469	5.008500	4.978663	5.004081	16.725573
4	2023-11-03	4.964172	5.000215	4.959903	4.996536	16.949241

### Exploratory Statistics

We perform t-tests to compare the mean values of the stock price metrics (e.g., close, adjusted close) on days with scientifically verifiable black swan anomalies and days without anomalies. This will help identify if and when there are significant differences in stock price.

```

import pandas as pd
from scipy.stats import ttest_ind
from datetime import datetime

#T-TEST FUNCTION
def perform_t_tests(start_date_anomaly, end_date_anomaly, start_date_before, end_date_before, start_date_after, end_date_after):
    # Filter the DataFrame for the specified date ranges: ANOMALY, BEFORE, AFTER
    selected_days_anomaly = logs_df[(logs_df['timestamp'] >= start_date_anomaly) & (logs_df['timestamp'] <= end_date_anomaly)]
    selected_days_before = logs_df[(logs_df['timestamp'] >= start_date_before) & (logs_df['timestamp'] <= end_date_before)]
    selected_days_after = logs_df[(logs_df['timestamp'] >= start_date_after) & (logs_df['timestamp'] <= end_date_after)]

    columns_to_compare = ['close', 'open', 'low', 'volume', 'high']
    significance_level=0.05

    # T-TESTS
    for column in columns_to_compare:
        t_statistic_before_after, p_value_before_after = ttest_ind(selected_days_before[column], selected_days_after[column], equal_var=False)
        t_statistic_anomaly_before, p_value_anomaly_before = ttest_ind(selected_days_anomaly[column], selected_days_before[column], equal_var=False)
        t_statistic_anomaly_after, p_value_anomaly_after = ttest_ind(selected_days_anomaly[column], selected_days_after[column], equal_var=False)

        # BEFORE VS AFTER
        if p_value_before_after < significance_level:
            print(f'T-Test for {column} - Before vs After:')
            print(f'T-Statistic: {t_statistic_before_after}')
            print(f'P-Value: {p_value_before_after}')
            print('\n')

        # ANOMALY VS BEFORE
        if p_value_anomaly_before < significance_level:
            print(f'T-Test for {column} - Anomaly vs Before:')
            print(f'T-Statistic: {t_statistic_anomaly_before}')
            print(f'P-Value: {p_value_anomaly_before}')
            print('\n')

        # ANOMALY VS AFTER
        if p_value_anomaly_after < significance_level:
            print(f'T-Test for {column} - Anomaly vs After:')
            print(f'T-Statistic: {t_statistic_anomaly_after}')
            print(f'P-Value: {p_value_anomaly_after}')
            print('\n')

#THIS CODE BOX CONTAINS 9/11 T-TEST

#4 Months ANOMALY 9/11 Financial Crisis
anomaly_1_start_date = '2001-09-11'
anomaly_1_end_date = '2002-01-11'

#4 Months BEFORE 9/11 Financial Crisis
b_anomaly_1_start_date = '2001-05-10'
b_anomaly_1_end_date = '2001-09-10'

#4 Months AFTER 9/11 Financial Crisis
a_anomaly_1_start_date = '2002-01-11'
a_anomaly_1_end_date = '2002-05-01'

perform_t_tests('2001-09-11', '2002-01-11', '2001-05-10', '2001-09-10', '2002-01-11', '2002-05-01' )

T-Test for close - Before vs After:
T-Statistic: 2.111784983046713
P-Value: 0.04524875041666402

T-Test for close - Anomaly vs After:
T-Statistic: 2.429314204473413
P-Value: 0.02111703762250402

T-Test for low - Before vs After:
T-Statistic: 2.1588812563206043
P-Value: 0.04144662691352722

T-Test for volume - Before vs After:
T-Statistic: -2.5162782857218877

```

P-Value: 0.017260649616976455

T-Test for high - Anomaly vs After:  
T-Statistic: 2.071689124225216  
P-Value: 0.04678480903365964

#THIS CODE-BOX CONTAINS 2008 T-TEST

#1 Year ANAMOLY 2008 Housing Financial Crisis  
anomaly\_2\_start\_date = '2008-01-01'  
anomaly\_2\_end\_date = '2008-12-31'

#1 Year BEFORE 2008 Financial Crisis

b\_anomaly\_2\_start\_date = '2007-01-01'  
b\_anomaly\_2\_end\_date = '2007-12-31'

#1 Year AFTER 9/11 Financial Crisis  
a\_anomaly\_2\_start\_date = '2009-01-01'  
a\_anomaly\_2\_end\_date = '2009-12-31'

perform\_t\_tests('2008-01-01', '2008-12-31', '2007-01-01', '2007-12-31', '2009-01-01', '2009-12-31')

T-Test for volume - Anomaly vs Before:  
T-Statistic: 2.703680566429757  
P-Value: 0.008039183516076407

T-Test for volume - Anomaly vs After:  
T-Statistic: 2.3893412999003156  
P-Value: 0.01874745320618078

T-Test for high - Anomaly vs Before:  
T-Statistic: 2.0613814883498924  
P-Value: 0.04267754786835433

#THIS CODE-BOX CONTAINS COVID-19 T-TEST

#2 Month ANAMOLY COVID-19 Financial Crisis  
anomaly\_3\_start\_date = '2020-03-01'  
anomaly\_3\_end\_date = '2020-04-31'

#2 Months BEFORE COVID-19 Financial Crisis  
b\_anomaly\_3\_start\_date = '2020-01-01'  
b\_anomaly\_3\_end\_date = '2020-02-29'

#2 Months AFTER COVID-19 Financial Crisis  
a\_anomaly\_3\_start\_date = '2020-05-01'  
a\_anomaly\_3\_end\_date = '2020-07-01'

perform\_t\_tests('2020-03-01', '2020-04-31', '2020-01-01', '2020-02-29', '2020-05-01', '2020-07-01')

T-Test for close - Anomaly vs Before:  
T-Statistic: -5.555126166757153  
P-Value: 0.00018585660045399157

T-Test for open - Anomaly vs Before:  
T-Statistic: -5.390561496627213  
P-Value: 0.00038636806876509366

T-Test for low - Anomaly vs Before:  
T-Statistic: -6.328028509682065  
P-Value: 6.838526834613872e-05

T-Test for high - Anomaly vs Before:  
T-Statistic: -5.878570899432568  
P-Value: 5.904096192407748e-05

## Machine Learning

We perform multiple regression to analyze the relationship between a single dependent (target) variable — volume — and several independent (predictor) variables, such as open, close, high and low stock price. We also created regressor trees in which our target variable took on continuous values instead of class labels in the leaves.

We varied all the parameters with the log transformation. We also performed our machine learning methods on all four of our independent variables and tried to take out additional ones to see how it would change our results. For example, we would look at high and low stock prices and omit open and close stock prices.

```
import matplotlib.pyplot as plt
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from statsmodels.stats.outliers_influence import variance_inflation_factor

def mult_lin_reg(df, features, target):
    # convert columns into numpy arrays
    X = df[features].values
    y = df[target].values

    # split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    #train linear regression model
    lin_model = LinearRegression().fit(X_train, y_train)

    # Calculate VIF for each feature
    df_features = pd.DataFrame(X, columns=features)
    vif_data = pd.DataFrame()
    vif_data["Variable"] = df_features.columns
    vif_data["VIF"] = [variance_inflation_factor(df_features.values, i) for i in range(df_features.shape[1])]

    print("Variance Inflation Factor:")
    print(vif_data)

    # generate predictions
    y_train_pred = lin_model.predict(X_train)
    y_test_pred = lin_model.predict(X_test)

    # evaluate scores (coefficient of determination, R^2)
    train_r2 = lin_model.score(X_train, y_train)
    test_r2 = lin_model.score(X_test, y_test)

    print(f"Training R-squared: {train_r2}")
    print(f"Testing R-squared: {test_r2}")

    # plot final multiple regression model
    plt.scatter(y_test, y_test_pred)
    plt.xlabel("Actual Values")
    plt.ylabel("Predicted Values")
    plt.title("Actual vs. Predicted Values of Testing Set")
    plt.show()


mult_lin_reg(logs_df, ["open", "high", "low", "close"], "volume")
```

Variance Inflation Factor:

Variable	VIF
0 open	120073.014719
1 high	139226.667768
2 low	119722.461206
3 close	130485.120974

Training R-squared: 0.5344901432537585  
Testing R-squared: 0.5195177560977121

Actual vs. Predicted Values of Testing Set




```
mult_lin_reg(logs_df, ["open", "close", "high"], "volume")
```

Variance Inflation Factor:

Variable	VIF
0 open	41768.493671
1 close	39768.307110
2 high	83995.761698

Training R-squared: 0.4790791402045186  
Testing R-squared: 0.49567243765317404

Actual vs. Predicted Values of Testing Set



```
mult_lin_reg(logs_df, ["open", "close", "low"], "volume")
```

Variance Inflation Factor:  
Variable VIF

0 open 35110.499663


1 close 40738.182948

2 low 72228.830026

Training R-squared: 0.47390902251660805

Testing R-squared: 0.4266113601708894

Actual vs. Predicted Values of Testing Set



```
mult_lin_reg(logs_df, ["open", "high", "low"], "volume")
```

Variance Inflation Factor:  
Variable VIF

0 open 64033.012674


1 high 43467.342640

2 low 36488.141864

Training R-squared: 0.5342192971911763

Testing R-squared: 0.5152884393953694

Actual vs. Predicted Values of Testing Set



```
mult_lin_reg(logs_df, ["close", "high", "low"], "volume")
```

Variance Inflation Factor:  
Variable VIF


0 close 69585.621921

1 high 40711.211284

2 low 41646.550433

Training R-squared: 0.5336692907921204

Testing R-squared: 0.5170137688249116



```
mult_lin_reg(logs_df, ["open", "close"], "volume")
```


Variance Inflation Factor:  
Variable VIF

0 open 19757.53683

1 close 19757.53683

Training R-squared: 0.2591800642165357

Testing R-squared: 0.21635312319498323



```
mult_lin_reg(logs_df, ["high", "low"], "volume")
```



Variance Inflation Factor:  
Variable VIF


0 high 23779.399401

1 low 23779.399401

Training R-squared: 0.5336512682659659

Testing R-squared: 0.5181592393363503

Actual vs. Predicted Values of Testing Set



```
mult_lin_reg(logs_df, ["open", "low"], "volume")
```

Variance Inflation Factor:  
Variable VIF


0 open 35030.128154

1 low 35030.128154

Training R-squared: 0.37468329081261353


Testing R-squared: 0.3938097885600982

Actual vs. Predicted Values of Testing Set



```
mult_lin_reg(logs_df, ["close", "high"], "volume")
```

Variance Inflation Factor:  
 Variable VIF  
 0 close 39732.085347  
 1 high 39732.085347  
 Training R-squared: 0.34955603303675586  
 Testing R-squared: 0.4030852708397965



With our multiple regression machine learning method, we notice that our models have high correlation but extremely high variance inflation factors. This means we are seeing multicollinearity, which is expected since our variables are related to each other — open, close, high and low prices rise and fall in similar patterns.

Therefore, we discovered that principal component analysis (PCA) could be incorporated to transform our original variables into a set of linearly uncorrelated variables. We created a function that adds PCA to our multiple regression in attempts to reduce multicollinearity. Hopefully, we can preserve as much information as possible with our modified variables.

```
def mult_lin_reg_with_pca(df, features, target, n_components=None):
    # Convert columns into numpy arrays
    X = df[features].values
    y = df[target].values

    # Split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Standardize the data
    scaler = StandardScaler()
    X_train_scaled = scaler.fit_transform(X_train)
    X_test_scaled = scaler.transform(X_test)

    # Apply PCA
    pca = PCA(n_components=n_components)
    X_train_pca = pca.fit_transform(X_train_scaled)
    X_test_pca = pca.transform(X_test_scaled)

    # Train linear regression model with PCA components
    lin_model_pca = LinearRegression().fit(X_train_pca, y_train)

    # Calculate VIF for each PCA component
    vif_data_pca = pd.DataFrame()
    vif_data_pca["Principal Component"] = range(1, pca.n_components_ + 1)
    vif_data_pca["VIF"] = [variance_inflation_factor(X_train_pca, i) for i in range(pca.n_components_)]

    print("Variance Inflation Factor (PCA Components):")
    print(vif_data_pca)

    # Generate predictions
    y_train_pred_pca = lin_model_pca.predict(X_train_pca)
    y_test_pred_pca = lin_model_pca.predict(X_test_pca)

    # Evaluate scores (coefficient of determination, R^2)
    train_r2_pca = lin_model_pca.score(X_train_pca, y_train)
    test_r2_pca = lin_model_pca.score(X_test_pca, y_test)
```

```

print(f"Training R-squared with PCA: {train_r2_pca}")
print(f"Testing R-squared with PCA: {test_r2_pca}")

# Plot final multiple regression model with PCA
plt.scatter(y_test, y_test_pred_pca)
plt.xlabel("Actual Values")
plt.ylabel("Predicted Values")
plt.title("Actual vs. Predicted Values of Testing Set with PCA")
plt.show()

```

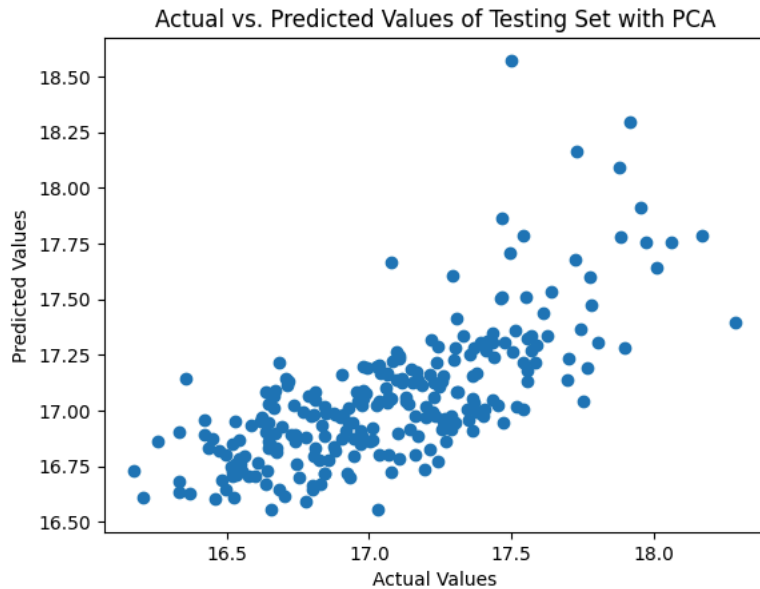
```
mult_lin_reg_with_pca(logs_df, ["open", "close", "high", "low"], "volume")
```

Variance Inflation Factor (PCA Components):

Principal Component	VIF
0	1 1.0
1	2 1.0
2	3 1.0
3	4 1.0

Training R-squared with PCA: 0.5344901432537585

Testing R-squared with PCA: 0.5195177560977086



```
mult_lin_reg_with_pca(logs_df, ["open", "close", "high"], "volume")
```


Variance Inflation Factor (PCA Components):

Principal Component	VIF
0	1 1.0
1	2 1.0
2	3 1.0

Training R-squared with PCA: 0.4790791402045186

Testing R-squared with PCA: 0.4956724376531755

Actual vs. Predicted Values of Testing Set with PCA



```
mult_lin_reg_with_pca(logs_df, ["open", "close", "low"], "volume")
```


Variance Inflation Factor (PCA Components):

Principal Component	VIF
0	1 1.0
1	2 1.0
2	3 1.0

Training R-squared with PCA: 0.4739090225166084

Testing R-squared with PCA: 0.42661136017088497

Actual vs. Predicted Values of Testing Set with PCA



```
mult_lin_reg_with_pca(logs_df, ["open", "high", "low"], "volume")
```

```
mult_lin_reg_with_pca(logs_df, ["close", "high", "low"], "volume")
```

```
mult_lin_reg_with_pca(logs_df, ["open", "close"], "volume")
```

```
mult_lin_reg_with_pca(logs_df, ["high", "low"], "volume")
```

We also created regressor trees.

Regressor trees can be used to build predictive models that learn patterns and relationships that help predict future stock prices.

Regressor trees also provide an intuitive way to assess the importance of different features in predicting stock prices. The interpretations of the trees help us understand which variables have the most significant impact on the model's predictions.

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.metrics import mean_squared_error

# create a function for regressor trees
def decision_tree_reg(logs_df, features, target):
    # Convert columns into numpy arrays
    X = logs_df[features].values
    y = logs_df[target].values

    # Split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Train Decision Tree Regressor
    tree_reg = DecisionTreeRegressor(random_state=42)
    tree_reg.fit(X_train, y_train)

    # Generate predictions
    y_train_pred = tree_reg.predict(X_train)
    y_test_pred = tree_reg.predict(X_test)

    # Evaluate model
    train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
    test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

    print("Train RMSE:", train_rmse)
    print("Test RMSE:", test_rmse)

    # Plot final model
    plt.scatter(y_test, y_test_pred)
    plt.xlabel("Actual Values")
    plt.ylabel("Predicted Values")
    plt.title("Actual vs. Predicted Values of Testing Set")
    plt.show()
```

We realized we needed to cross-validate our regressor trees to better estimate our model's performance, ensure it can react well to unseen testing data and detect overfitting.

```
from sklearn.tree import DecisionTreeRegressor, plot_tree
from sklearn.model_selection import cross_val_score, KFold
from sklearn.metrics import mean_squared_error

def decision_tree_reg_with_r2(logs_df, features, target, max_depth=10, min_samples_leaf=50):
    # Convert columns into numpy arrays
    X = logs_df[features].values
    y = logs_df[target].values

    # Split data into training and testing sets
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

    # Train Decision Tree Regressor with limited depth and minimum samples per leaf
    tree_reg = DecisionTreeRegressor(min_samples_leaf=min_samples_leaf, random_state=42)
    tree_reg.fit(X_train, y_train)

    # Generate predictions
    y_train_pred = tree_reg.predict(X_train)
    y_test_pred = tree_reg.predict(X_test)

    # Evaluate model
    train_rmse = np.sqrt(mean_squared_error(y_train, y_train_pred))
    test_rmse = np.sqrt(mean_squared_error(y_test, y_test_pred))

    print("Train RMSE:", train_rmse)
    print("Test RMSE:", test_rmse)

    # Perform cross-validation with R-squared as the scoring metric
    kfold = KFold(n_splits=5, shuffle=True, random_state=42)
    r2_scores = cross_val_score(tree_reg, X, y, cv=kfold, scoring='r2')

    print("Cross-Validation R-squared Scores:", r2_scores)
    print("Average R-squared:", np.mean(r2_scores))

    # Plot final model
    plt.scatter(y_test, y_test_pred)
    plt.xlabel("Actual Values")
    plt.ylabel("Predicted Values")
    plt.title("Actual vs. Predicted Values of Testing Set")
    plt.show()

    # Visualize the Decision Tree
    plt.figure(figsize=(30, 8))
    plot_tree(tree_reg, filled=True, feature_names=features, rounded=True, fontsize=10)
    plt.show()

decision_tree_reg(logs_df, ["open", "close", "high", "low"], "volume")
```



```
decision_tree_reg(logs_df, ["open", "close", "high"], "volume")
```

```
decision_tree_reg(logs_df, ["open", "close", "low"], "volume")
```

```
decision_tree_reg(logs_df, ["open", "high", "low"], "volume")
```

```
decision_tree_reg(logs_df, ["close", "high", "low"], "volume")
```

```
decision_tree_reg(logs_df, ["open", "close"], "volume")
```

```
decision_tree_reg(logs_df, ["high", "low"], "volume")
```

```
decision_tree_reg(logs_df, ["open", "low"], "volume")
```

```
decision_tree_reg(logs_df, ["close", "high"], "volume")
```



## Visualizations

### Line graph

We create line graphs showing the progression of stock prices. We generated a line plot comparing the the average values of "high" and "low" stock prices from 1991 to 2023, with emphasis on the speicified anomaly dates which are easily identified and when the line graphs become their lowest point. The resulting line plot display a visual representation of how low the stock dropped during the specified anomaly times.

```
import matplotlib.dates as mdates
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

def plot_stocks(df, col1, col2, interval_val, start_date=None, end_date=None):
    df["timestamp"] = pd.to_datetime(df["timestamp"])

    if start_date is not None and end_date is not None:
        start_date = pd.to_datetime(start_date)
        end_date = pd.to_datetime(end_date)
        df = df[(df["timestamp"] >= start_date) & (df["timestamp"] <= end_date)]

    df = df.sort_values(by='timestamp')

    # Plot
    plt.figure(figsize=(15, 10))
    plt.gca().xaxis.set_major_formatter(mdates.DateFormatter("%Y-%m-%d"))
    plt.gca().xaxis.set_major_locator(mdates.DayLocator(interval=interval_val))

    plt.plot(df["timestamp"], df[col1], label=str(col1).capitalize())
    plt.plot(df["timestamp"], df[col2], label=str(col2).capitalize())

    plt.xlabel("Timestamp")
    plt.ylabel("Price (in USD)")
    plt.legend()
    plt.gcf().autofmt_xdate()
    plt.title(f"Difference in {str(col1).capitalize()} Stock Prices and {str(col2).capitalize()} Stock Prices over {start_date}")
    plt.show()
```